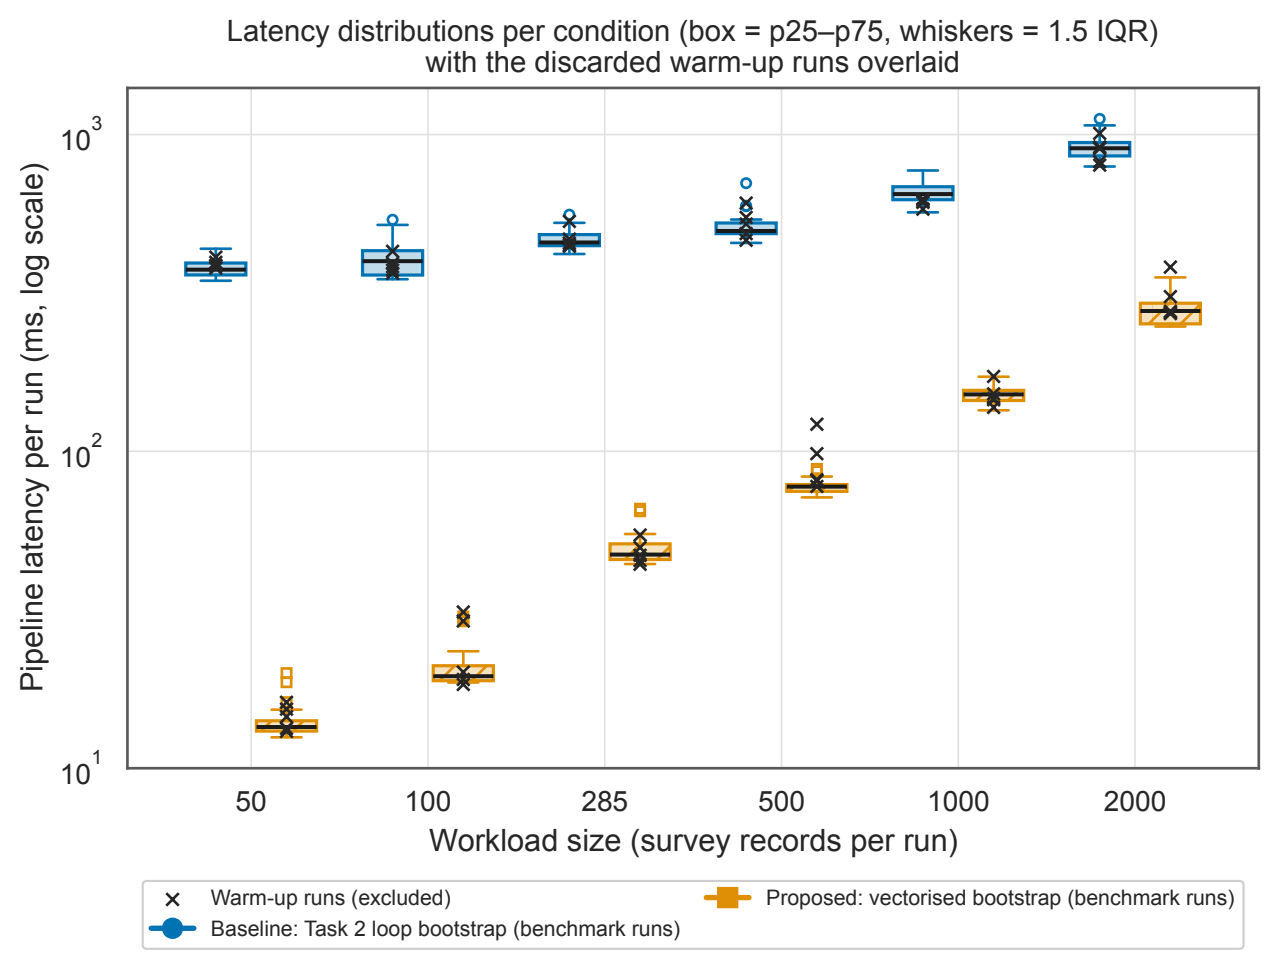
Latency distributions per condition (box = p25–p75, whiskers = 1.5 IQR)
with the discarded warm-up runs overlaid
Pipeline latency per run (ms, log scale)
10^3
10^2
10^1
50
100
285
500
1000
2000
Workload size (survey records per run)
Warm-up runs (excluded)
Baseline: Task 2 loop bootstrap (benchmark runs)
Proposed: vectorised bootstrap (benchmark runs)